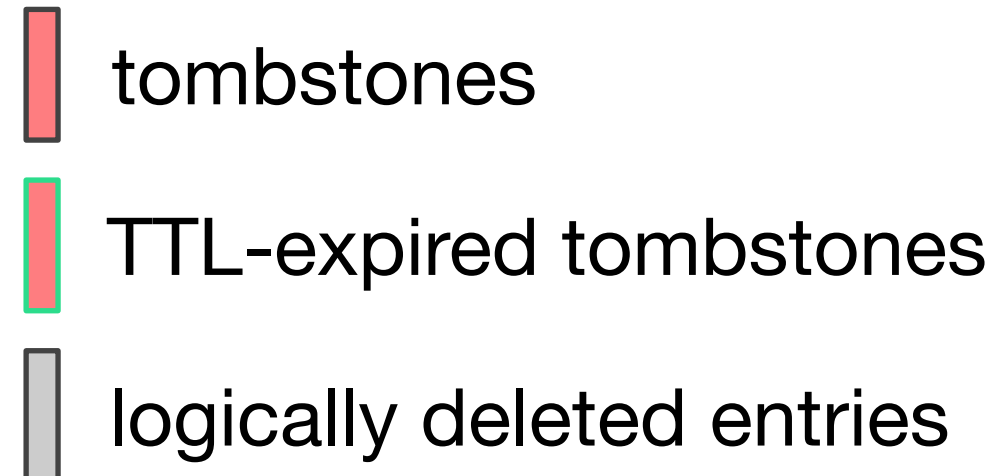


## setting environment

after every flush, perform the following check


```
//Update d[i] if a new level is added to the tree
// T = size ratio; level_count = levels in tree
// D_th = delete persistence threshold
double x = D_th * T / pow ( T, level_count - 1 );
if ( new_level_added () ) {
  for ( int i = 0; i < level_count; ++i ) {
    d[i] = x * pow ( T, i );
    if ( i > 0 )
      d[i] += d[i-1];
  }
}
```



## FADE: enforcing a finite bound for delete persistence latency

initial state\*


Levels 1, 2, and 3  
all have files that  
contain tombstone



Choose level that  
has files with  
expired TTL


selecting files for compaction

Levels 2 and 3 both  
have files with  
expired TTL



Chose the file that  
has the highest  
tombstone count


Level 2 has the file  
with highest  
tombstone count



Initiate compaction  
with chosen file and  
update age of files


performing compaction

After compaction,  
Level 3 has 2 files  
with tombstones



Initiate compaction  
with the file that has  
expired TTL

All files with expired  
TTL are compacted —  
repeat routine



This ensures a finite  
upper-bound for delete  
persistence latency

\* timers are shown only for files containing tombstones